Iterative Knowledge Distillation

Sharan Yalburgi

July 15, 2019

1 Abstract

Deep (Convolutional) Neural Networks have recently shown remarkable success in many real-world visual recognition tasks. These networks are defined with very deep structure (hidden layers and neurons) and numerous model parameters (weights). This makes the existing deep neural network models computationally expensive and memory intensive. Such expensive models find very limited application when it comes to their deployment in devices with lower memory resources or applications with strict latency requirements. Therefore, a straightforward idea is to compress these deep models and accelerate the training process without significantly degrading the original model performance. In this paper, we develop an approach of recursively generating student models (compressed models) from the teacher models (original model and subsequently the compressed models). Each student model is trained with a composite loss function consisting of three different functional terms: KL-divergence between output logits of student and the teacher model, the cross-entropy loss with true labels, and additionally a hint from the teacher model. We also demonstrate the student model building as a search procedure. The recursive building of student (compressed) models stops when there is a significant drop in the performance. We evaluate our model compression approach for compressing Deep Convolutional Neural Networks (CNN) for two well-known benchmark datasets such as MNIST digits dataset and CIFAR-10 dataset. Experimental results demonstrate that our proposed approach can compress original Deep-CNN model with x% reduction in memory footprint with y% drop in performance as compared with the original model. These results are comparable/better to/than the state-of-the-art model compression methods.

2 Introduction

Problem Statement: What is the problem? Why is it important to solve it? How do we intend to solve it?

Deep Neural Networks have achieved remarkable results in the recent past. They have been applied in various fields including medical diagnosis, finance, drug discovery, speech recognition and space exploration to name a few.

However, they often have hundreds of millions of parameters. Hence take several days to train even on high performance GPUs. This not just slows down the initial training procedure but also is burden while eventually deploying on real time system where predictions are often required to almost instantaneous with little resource..

As we entertain larger and larger neural networks to model our data, it becomes more important to consider its storage and computational requirements. Especially in the context of real time systems we need models to perform with less resource both in terms of computation and storage. Another aspect which is of importance in remote embedded systems is power consumption.

We will restrict or investigation to storage and computation in this report.

3 Model Compression

Model compression is the process of modifying or using a trained model to create a new model which generally takes lesser computation and/or storage and/or power but retains comparable performance.

There are several existing techniques for model compression. Some of them include parameter pruning, low-rank factorization, knowledge distillation to name a few.

We focus on the idea of model compression by knowledge distillation in this report.

3.1 Knowledge Distillation

Knowledge Distillation is the idea of complex model transferring its learnt knowledge to a less complex model. We aim to match the original model's performance in the learnt model. Throughout this report we use the terminology that *parent model* is teaching the *child model*. We use the term *ancestors* to refer to any and all models which are used to train a specific model.

The idea Knowledge Transfer was first proposed in Bucilă, Caruana, and Niculescu-Mizil, 2006. A more recent work first applied it to deep neural networks.Ba and Caruana, 2014

3.2 Iterative Knowledge Distillation

The key idea of iterative knowledge distillation is to gradually reduce the model size by using a parent model(s) to train a child mode(s). Throughout this process the of learning and re-learning we keep the algorithm exposed to the data. We have seen that this at times increases the performance with child model(s).

We propose several variations of Iterative Knowledge Distillation.

Parent → Child

This is simplest version of Iterative Knowledge Distillation where a single parent teaches a single child. This process could be repeated until the drop in the performance is within reasonable amounts.

Parent → Children

In this variation we propose in ensemble learning approach where at each learning phase, we train an ensemble of children using a single parent. The children could be of same architecture with different initialisations or could have different architectures. We could select the best performing child at each step before performing the next iteration of Knowledge Distillation or we could grow a tree of Parent-Child models.

Ancestors → Child

We propose that several expert Parent models could be used to train a single Child model. The parent could have different architectures catering to different aspects of the data.

Ancestors → Children

This the a development over the previous approach where several expert Parent models could help train a several Child models. We could select the best performing child at each step before performing the next iteration of Knowledge Distillation or we could grow a tree of Ancestors-Children models.

4 Empirical Evaluation

4.1 Data

We consider two of the most popular datasets in vision. MNIST and CIFAR-10

MNIST

MNIST(Deng, 2012) is a dataset of handwritten digits. It hence has 10 classes. It has training set of 60,000 digits and test set of 10,000 digits. All digits are greyscale images of dimensions 28x28. The digits have been size-normalized and centered in a fixed-size image. Each class has been equally represented in the train and test set.

CIFAR-10

CIFAR-10(Krizhevsky and Hinton, 2010) is a dataset of 10 classes. It has training set of 50,000 digits and test set of 10,000 digits. All digits are RGM images of dimensions 32x32. Each class has been equally represented in the train and test set. The classes are completely mutually exclusive. The dataset contains common objects like animals and automobiles.

4.2 Model Compression for MNIST and CIFAR-10

To demonstrate model compression, we compress models trained on MNIST and CIFAR-10 to the smallest possible size. We consider VGG-16 and

VGG-19 models as our initial trained models. We
two different compression strategies.

Reduce number of Blocks

Here we reduce the number of blocks in the VGG-
16 and VGG-19 network in a gradual manner.

Reduce Block Size

We reduce the complexity blocks in the VGG-16
and VGG-19 network in a gradual manner.

4.3 Loss Function

$LOSS = KLDivergence(ParentModel, ChildModel) + CrossEntropyLoss(Y, \hat{Y})$

5 Results

$LOSS = KLDivergence(ParentModel, ChildModel) + CrossEntropyLoss(Y, \hat{Y})$

6 Conclusion

References

Ba, Jimmy and Rich Caruana (2014). "Do deep nets
really need to be deep?" In: *Advances in neural
information processing systems*, pp. 2654–2662
(cit. on p. 2).

Buciluă, Cristian, Rich Caruana, and Alexandru
Niculescu-Mizil (2006). "Model compression". In:
*Proceedings of the 12th ACM SIGKDD inter-
national conference on Knowledge discovery and
data mining*. ACM, pp. 535–541 (cit. on p. 2).

Deng, Li (2012). "The MNIST database of hand-
written digit images for machine learning re-
search [best of the web]". In: *IEEE Signal Pro-
cessing Magazine* 29.6, pp. 141–142 (cit. on p. 2).

Krizhevsky, Alex and Geoff Hinton (2010). "Convo-
lutional deep belief networks on cifar-10". In: *Un-
published manuscript* 40.7, pp. 1–9 (cit. on p. 2).